

DSCI6672 – Spring 2020

Midterm Project Cybersecurity Report

Nitin Pathania

Overview and Tasks

The requirement of the midterm project is to train a deep neural network, deploy the model through Amazon Sage Maker service and create a python script that takes in a PE file(Putty) as a parameter and classify as one of the following binary classes: Benign or Malicious. The report is divided into three tasks based the complications and approach to creating a solution for each one of them.

Task 1: Preprocessing

As the name of this task, the data must be preprocessed and vectorized to be fed into a deep neural network. We have used LIEF project to convert data to vectors easily. Initially I was not aware of the library and was trying to figure out how to do this but after becoming aware of the library, I immediately scrapped my vectorized data and installed their EMBER library.

The notebook was created on Amazon Sagemaker but still the student tier we had access does not allow us to create big instances and there was a limited access to amazon web services which led to a major constraint such as: Overloading memory crash and session crashes. Due to this, I decided to download the data to the runtime session every time the session timed out or crashed. Therefore, The normal sagemaker instances comes with 5 GB disk space. This dataset required a lot more so I had to tweak the space several time and yes I wasted a lot of time there. I was used to host the data files in google collabration but aws sagemaker made my life easier in many ways. The files are referenced through numpy's memmap, the data does not have to be loaded into the RAM at once; however, this would lead to slower processing of data to a

machine learning model. It took a while to download data first like 1 hour then creating vectors took another hour or so. It took me 5 days to reach this step.

After constructing the full training data, full test data and their labels, the training data had to be separated from the valid classes and the unlabeled classes. however, it loaded all the data into the RAM, it was also a longer process. It should be noted that by being able to load the full dataset into the RAM, the model's training takes less time.

The last preprocessing step for the data is to standardize or normalize it. This is needed to allow the neural network to converge, if not at least faster. The chosen normalizer was Standard Scaler from the Sklearn library. The data was too large to fit to the scaler instance at once, `partial_fit` was used on partitions of the training data. Finally, by creating two scaler instances for test and train data, we transform the data and rewrite it in memory . The data was ready to be given to a neural network.

The considered architectures for the neural network were between two general concepts: A model with Conv1D as its first layer after input layer or a typical dense layer. It should be noted that because of the upgrade in Tensorflow between 1.x and 2.x, there is some required reshaping if using a Conv1D layer, Eg. Data in the format of $(v[1], v[2], v[3])$ where there is 2 axes, (an array of an array of numbers) to 3 axes in the format $([v[1]], [v[2]], [v[3]])$ (an array of tuples.) The selected architecture is as followed:

Layer (type)	Output Shape	Param #
dropout (Dropout)	(None, 1, 2381)	0
dense (Dense)	(None, 1, 1500)	3573000
dropout_1 (Dropout)	(None, 1, 1500)	0
dense_1 (Dense)	(None, 1, 1)	1501
Total params: 3,574,501		
Trainable params: 3,574,501		
Non-trainable params: 0		

It is recommended to add dropout layers to produce better generalization. The decision for the number of nodes in the dense layer is the size of approximately 2/3 of the number of features: 2381. The choice of activation functions were the following: relu for hidden layers, sigmoid for output layer and Adam as the optimizer with the set learning rate of .01. There is a L1 regularizer set to the dense layer at .01 which is to prevent overfitting. The selected performance metrics were accuracy, AUC, and precision because false-alarm rate and detection rate are import to this domain. The wide architecture was recommended from the book, Deep Learning with Python by François Chollet.

Task 2: Training

For training, it was recommended to do small batches. For example, the model that showed the dataset did batches of 256 and MalConv that came with EMBER did batches of 100. Time variated depending on the RAM chip, but within an hour, the model finished through one epoch, providing a promising 87-90% accuracy, 96% AUC, and 96% precision on training data and 98%,99%,98% on validation data respectively.

After saving the first preliminary round, the model was set to run for 30 epochs to see if it'd produce better results. The model ran for all 30 epochs but despite longer training, the model seemed to get progressively worse as all the performance metrics continued to dip. When this happened, there were some alterations to the learning rate and batch size, but neither benefitted the model in this case. Those models were scrapped for the best performing model that occurred after the first epoch.

The final evaluation was done on the test set, receiving 97.5%,98.5%,97.7% in accuracy, auc, and precision respectively. Additional metrics like f-score, 97.2%, and the confusion matrix were also calculated.

Task 2: Deploy of Model on Cloud

The service used to deploy the model was Amazon's Web Services's SageMaker. Caution: Tensorflow 2.x is currently the most recent distribution of Tensorflow. AWS supports Tensorflow 2.0 for training and deploying housed TensorFlow Serving models from TensorFlow Estimators; however, to import an outside trained model, TensorFlowModel method is used to convert a TensorFlow Model into an estimator. The caveat is the framework that supports TensorFlow 1.x models, like 1.12 and 1.6 according to the tutorial and method documentation's default framework TensorFlow version. If one's intention is to deploy an outside model through SageMaker's Endpoint, it would be wise to assure their TensorFlow version is compatible with the TensorFlowModel method. In the tutorial resource on the

deployment of a model through SageMaker, TensorFlow1.12 is compatible. It tried the code on both AWS and google Collab now. Google Collaboratory's current 1.x version is 1.5, but I could not deploy a functional model hosted by the Endpoint and Sagemaker process is quite difficult due to tensorflow and keras versioning issues. Therefore, downgrading to Tensorflow 1.12, recreating the architecture and loading the weights to resave the model was necessary. After getting the model to successfully import, the rest of the tutorial provided was sufficient. Depending the Amazon Web Service account; for example, an AWS educate account stores their keys and password before launching the AWS console; however, cannot be found elsewhere. A normal user AWS account can create their passwords and keys through the IAM, Identity Access Management.

The creation of the endpoint took around 6 minutes. When invoking the endpoint from within the SageMaker notebook, response time was optimal—less than half a second.

Task 3: Create a Client

I launched a virtual machine through AWS through EC2 service, specifically the distribution Windows to create an environment to write the code. This was an optimal testing environment for packages. Only one necessary file was bundled and that was the pickle scaler file. Normally, a client shouldn't have this file; instead, this file should be hosted with a Lambda function set to scale incoming data but said service was restricted in student account and chargeable in my personal account. The python script was written through the nano text editor and majority of the imports are already on the client except for EMBER. The python script installs EMBER if it is not installed. Putty.exe was used as the deployment sample and returned

benign. Note: The code included for the client side has the sensitive keys and passwords removed. The skeleton of the code was provided by the EMBER GitHub repository and adjusted for this model. The response time for invoking the endpoint and returning data was around 7 seconds. That is incredibly long for Malware detection and therefore should not be recommended for general use. This encourages the concept of having light models distributed on the client for more optimal speeds.

Conclusion

Most complications and delays of this project were due to version incompatibilities and limitations in the service's free tier than of the objective themselves. This project should be forward to implement after knowing some of these complications. I spent time learning what Lambda and other services of AWS because many resources used them; however, due to account restrictions, could not use them. SO I created AWS free tier account to learn them and implement them. I believe I have better understanding why all deployed systems are practically legacy, products would be too unreliable otherwise. I also found passion of listing package versions ahead of time, say for a tutorial, because it'd make it easier to read through documentations. Most of the packages were too big to download.

Bibliography

H. Anderson and P. Roth, "EMBER: An Open Dataset for Training Static PE Malware Machine Learning Models", in ArXiv e-prints. Apr. 2018

Chollet, François. *Deep Learning with Python* . : Manning, 2017.

Youtube and Google.